

# tutorial\_ebm\_CIFAR10

June 5, 2023

## 1 Tutorial on iterative image denoising with Langevine dynamics.

```
[ ]: # Author: Róisín Luo  
# Ref: https://uvadlc-notebooks.readthedocs.io/en/latest/tutorial_notebooks/  
↪tutorial8/Deep_Energy_Models.html
```

## 2 Importing libraries

```
[143]: import sys  
import os  
import random  
import math  
  
import numpy as np  
  
import matplotlib  
import matplotlib.pyplot as plt  
from matplotlib import cm  
from matplotlib import patches  
#from matplotlib.colors import to_rgb  
#from mpl_toolkits.mplot3d.axes3d import Axes3D  
#from mpl_toolkits.mplot3d import proj3d  
  
from tqdm import tqdm  
  
#import csv  
#import pandas as pd  
  
import torch  
import torch.nn as nn  
import torch.nn.functional as F  
  
import torchvision  
from torchvision import transforms  
from torchvision.utils import make_grid
```

```

from PIL import Image

from atomicwrites import atomic_write #we must guarantee the atomicity of
    write operation.
import pickle

%matplotlib inline

```

### 3 Find GPU

```

[144]: def find_gpu():

    device = torch.device('cpu')

    if torch.cuda.is_available():
        print('# of CUDA devices: ', torch.cuda.device_count())

        for i in range(torch.cuda.device_count()):
            print()
            print(torch.cuda.get_device_name(i))
            print('CUDA memory Usage:')
            print('Allocated:', round(torch.cuda.memory_allocated(i)/
    ↪1024**3,1), 'GB')
            print('Cached:   ', round(torch.cuda.memory_reserved(i)/1024**3,1),
    ↪'GB')

        device = torch.device('cuda')

        # MacOS
        elif hasattr(torch, "backends") and \
            hasattr(torch.backends, "mps") and \
            torch.backends.mps.is_available():

            device = torch.device('mps')

        print()
        print("Detected GPU device is: ", device)

        return device

```

```

[145]: device = find_gpu()
device

```

Detected GPU device is: mps

```
[145]: device(type='mps')
```

## 4 Load dataset

```
[150]: def plot_images(images, num_cols = 8, num_rows = 2):
```

```
    assert len(images) > num_cols * num_rows

    images = images[0: num_cols * num_rows]

    images = (images + 1) / 2

    #create grid with size by nrow x ncol
    img_grid = make_grid(tensor = images,
                          nrow = num_cols,
                          padding = 1,
                          pad_value = 1)

    #conver to channel last
    img_grid = img_grid.permute(1,2,0)

    plt.figure(figsize=(16, 16))
    plt.imshow(img_grid)
    plt.axis('off')
    plt.show()
```

```
[151]: transform = transforms.Compose([transforms.ToTensor(),
                                       transforms.Normalize((0.5,), (0.5,)) #to [-1,+1]
                                       ])


    dataset = torchvision.datasets.MNIST(root="datasets",
                                         train=True,
                                         transform=transform,
                                         download=True)

    dataloader = torch.utils.data.DataLoader(dataset,
                                              batch_size=128,
                                              shuffle=True,
                                              #drop_last=True,
                                              #num_workers=4,
                                              #pin_memory=True
                                              )
```

```
[152]: x,y = next(iter(dataloader))
        x.shape
```

```
[152]: torch.Size([128, 1, 28, 28])
```

```
[153]: plot_images(x)
```



## 5 Build energy model (simple)

```
[154]: #we hope the energy model is smooth.
class Swish(nn.Module):
    def forward(self, x):
        return x * torch.sigmoid(x)

#As standard practice/convention, the outputs represent the negative energy.
class EBMSimple(nn.Module):

    def __init__(self, **kwargs):
        super().__init__()

        self.net = nn.Sequential(
            nn.Conv2d(1, 32, kernel_size=5, stride=2, padding=4),
            Swish(),

            nn.Conv2d(32, 64, kernel_size=3, stride=2, padding=1),
            Swish(),

            nn.Conv2d(64, 32, kernel_size=3, stride=2, padding=1),
            Swish(),

            nn.Conv2d(32, 32, kernel_size=3, stride=2, padding=1),
            Swish(),

            nn.Flatten(),
            nn.Linear(128, 100),
            Swish(),
```

```

        nn.Linear(100, 1)
    )

    def forward(self, x):
        x = self.net(x)
        return x

```

```

[155]: model = EBMSimple()
print(model)

x = torch.rand(10, 1, 28, 28)
x = model(x)

print(x.shape)
print(x)

```

```

EBMSimple(
  (net): Sequential(
    (0): Conv2d(1, 32, kernel_size=(5, 5), stride=(2, 2), padding=(4, 4))
    (1): Swish()
    (2): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    (3): Swish()
    (4): Conv2d(64, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    (5): Swish()
    (6): Conv2d(32, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    (7): Swish()
    (8): Flatten(start_dim=1, end_dim=-1)
    (9): Linear(in_features=128, out_features=100, bias=True)
    (10): Swish()
    (11): Linear(in_features=100, out_features=1, bias=True)
  )
)
torch.Size([10, 1])
tensor([[0.0904],
        [0.0904],
        [0.0907],
        [0.0905],
        [0.0906],
        [0.0905],
        [0.0906],
        [0.0905],
        [0.0904],
        [0.0902]], grad_fn=<AddmmBackward0>)

```

## 6 Build energy model (resnet based)

```
[156]: #we hope the energy model is smooth.
class Swish(nn.Module):
    def forward(self, x):
        return x * torch.sigmoid(x)

#As standard practice/convention, the outputs represent the negative energy.
class EBMResNet(nn.Module):

    def __init__(self, **kwargs):
        super().__init__()

        self.decoder = torchvision.models.resnet18(pretrained = True)
        self.decoder.fc = nn.Identity()

        self.adapter = nn.Sequential(
            nn.Linear(512, 100),
            Swish(),

            nn.Linear(100, 1),

        )

    def forward(self, x):
        #self.decoder.eval()
        x = self.decoder(x)
        x = self.adapter(x)
        return x
```

```
[9]: #model = EBMResNet()
#x = torch.rand(10, 3, 32, 32)
#x = model(x)
#
#print(x.shape)
```

## 7 Sampler for training and generating

```
[157]: class EBMSampler:
        def __init__(self,
                     #model,          #EBM
                     image_size,    #image size
                     sample_size,   #sampling batch size.
                     buffer_len = 8192):
```

```

        #self.model = model
        self.buffer_len = buffer_len
        self.image_size = image_size
        self.sample_size = sample_size
        self.sample_shape = (1,) + image_size

        self.buffer = [torch.rand((1, ) + image_size) * 2 - 1] * self.buffer_len

#sample a batch for training.
def train_sample(self,
                model,
                num_new = None,
                iters = 60, step_size = 10, device = torch.device("cpu")):

    if num_new is None:
        num_new = np.random.binomial(self.sample_size, 0.05)

    if num_new == 0:
        num_new = 1

    num_old = self.sample_size - num_new

    new_images = torch.rand((num_new,) + self.image_size) * 2 - 1

    old_images = random.choices(self.buffer, k = num_old)
    old_images = torch.cat(old_images, dim = 0)

    images = torch.cat([new_images, old_images], dim = 0)
    images = images.detach().to(device)

    # Perform MCMC sampling
    images = self.langevine_dynamics_sample(model,
                                            images,
                                            iters = iters,
                                            step_size = step_size,
                                            device = device)

    self.buffer = list(images.to(torch.device("cpu")).chunk(self.
↪sample_size, dim=0)) + self.buffer
    self.buffer = self.buffer[:self.buffer_len]

    return images

#MCMC sampling with Langevine dynamics.
def langevine_dynamics_sample(self,
                            model,
                            images,

```

```

        iters = 60,
        step_size = 10,
        return_images_per_step = False,
        device = torch.device("cpu")):

model.to(device)
images = images.to(device)

images.requires_grad = True

noise = torch.randn(images.shape, device = images.device)

images_per_step = []

for _ in range(iters):

    noise.normal_(0, 0.005)
    images.data.add_(noise.data)
    images.data.clamp_(min=-1.0, max=1.0)

    outputs = -model(images)
    outputs = outputs.sum()

    #compute gradients wrt images.
    outputs.backward()

    #gradient clamp.
    images.grad.data.clamp_(-0.03, 0.03)

    # Towards the direction with low-energy
    images.data.add_(-step_size * images.grad.data)

    images.grad.detach_()
    images.grad.zero_()
    images.data.clamp_(min=-1.0, max=1.0)

    if return_images_per_step:
        images_per_step.append(images.clone().detach())

if return_images_per_step:
    return torch.stack(images_per_step, dim=0)
else:
    return images

```



```
[158]: model = EBMSimple()
sampler = EBMSampler(image_size = (1, 28, 28), sample_size = 10)

model.eval()
for p in model.parameters():
    p.requires_grad = False

images = sampler.train_sample(model = model, device = device)
images.shape
```

```
[158]: torch.Size([10, 1, 28, 28])
```

## 8 Training energy-based model for one epoch

```
[159]: %matplotlib inline

#from IPython.display import display, clear_output
from IPython import display

def EBM_train_one_epoch(
    model,
    device,
    dataloader,
    optimizer,
    epoch,
    sampler,
    sampler_iters = 60,
    sampler_step_size = 10,
    alpha = 0.2, #regularization weight.
    max_batches = None):

    # Enable gradient computing
    model.to(device)
    model.train()

    if max_batches is None:
        max_batches = len(dataloader)

    #averaged loss in current epoch.
    epoch_loss = 0.0
    total_loss_ = 0.0

    #contrastive divergence
    epoch_cdiv = 0.0
    total_cdiv_ = 0.0
```

```

#regularization
epoch_reg = 0.0
total_reg_ = 0.0

for batch_idx, batch in enumerate(dataloader, 1):

    real_images, _ = batch

    real_images = real_images.to(device)

    #simple data augmentation.
    noise = torch.randn_like(real_images) * 0.005
    #noise = torch.randn_like(real_images) * math.sqrt(eta)

    real_images.add_(noise).clamp_(min=-1.0, max=1.0)

    fake_images = sampler.train_sample(model = model,
                                       iters = sampler_iters,
                                       step_size = sampler_step_size,
                                       device = device)

    images = torch.cat([real_images, fake_images], dim = 0)

    optimizer.zero_grad()

    #model outputs negative energy.
    outputs = -model(images)

    real_outputs, fake_outputs = outputs.chunk(2, dim = 0)

    #computing batch loss.

    #we hope the minimum energy close to zero.
    batch_reg = alpha * (real_outputs**2 + fake_outputs**2).mean()
    batch_cdiv = real_outputs.mean() - fake_outputs.mean()
    batch_loss = batch_reg + batch_cdiv

    #computing gradients
    batch_loss.backward()

    #update weights.
    optimizer.step()

    #update statistics
    total_reg_ += batch_reg
    total_cdiv_ += batch_cdiv

```

```

total_loss_ += batch_loss

epoch_reg = total_reg_ / batch_idx
epoch_cdiv = total_cdiv_ / batch_idx
epoch_loss = total_loss_ / batch_idx

#Updating training displays.
display.clear_output(wait=True)

display.display('Epoch {} [{} / {}] {:.0f}%'.format(
    epoch, batch_idx,
    len(dataloader),
    100. * (batch_idx / len(dataloader))))

display.display('* regularization batch {:.6f} epoch {:.6f}'.format(
    batch_reg, epoch_reg))

display.display('* contrastive divergence batch {:.6f} epoch {:.6f}'.
    ↪format(
        batch_cdiv, epoch_cdiv))

display.display('* loss batch {:.6f} epoch {:.6f}'.format(
    batch_loss, epoch_loss))

if batch_idx > max_batches:
    break

return epoch_reg, epoch_cdiv, epoch_loss

```

```

[160]: device = find_gpu()
       print(device)

```

Detected GPU device is: mps  
mps

```

[161]: batch_size = 128
       model = EBMSimple()
       sampler = EBMSampler(image_size = (1, 28, 28), sample_size = batch_size)

```

```

[162]: dataloader = torch.utils.data.DataLoader(dataset,
        batch_size=batch_size,
        shuffle=True,
        #drop_last=True,
        #num_workers=4,
        #pin_memory=True

```

```
)
```

```
[163]: #for Adam the learning rate
#typically less than 0.001 for stabability.
learning_rate = 0.001

optimizer = torch.optim.Adam(
    model.parameters(),
    lr = learning_rate,
    betas = (0.99, 0.999),
    #momentum = 0.9,
    #weight_decay = 5e-4
)

scheduler = torch.optim.lr_scheduler.StepLR(optimizer, 1, gamma=0.97) #␣
    ↪Exponential decay over epochs
```

```
[61]: epoch_reg, epoch_cdiv, epoch_loss = EBM_train_one_epoch(
    model,
    device,
    dataloader,
    optimizer,
    epoch = 1,
    sampler = sampler,
    sampler_iters = 60,
    sampler_step_size = 10,
    alpha = 0.2,
    max_batches = None)
```

```
'Epoch 1 [469/469 (100%)]'
```

```
'* regularization batch 0.000214 epoch 0.012221'
```

```
'* contrastive divergence batch 0.010302 epoch -0.016724'
```

```
'* loss batch 0.010516 epoch -0.004503'
```

## 9 Save and load model from disk

```
[164]: import os

def save_model(model, model_path):

    save_path = os.path.normpath(os.path.dirname(model_path)).rstrip(os.path.
    ↪sep)

    if not os.path.exists(save_path):
```

```
os.makedirs(save_path)

print("Save model weights to: ", model_path)
torch.save(model.state_dict(), model_path)
```

```
[165]: save_model(model, "models/ebm_simple.pth")
```

Save model weights to: models/ebm\_simple.pth

```
[166]: def load_model(model_path, device):
        model = EBMSimple()

        if os.path.exists(model_path):
            #re-loading
            model.load_state_dict(torch.load(model_path, map_location = device))
            print("Loaded model weights from: ", model_path)
        else:
            print("Model weights not found.")

        return model
```

```
[167]: model = load_model("models/ebm_simple.pth", device)
```

Loaded model weights from: models/ebm\_simple.pth

## 10 Generate samples

```
[168]: def generate_samples(model,
                            sampler,
                            iters = 60,
                            step_size = 10,
                            num_samples = 16,
                            image_size = (1, 28, 28),
                            device = torch.device("cpu")):

    model.eval()
    model.to(device)

    for p in model.parameters():
        p.requires_grad = False

    images = torch.rand((16,) + image_size) * 2 - 1

    images = sampler.langevine_dynamics_sample(
        model,
        images,
        iters = iters,
```

```

        step_size = step_size,
        return_images_per_step = False,
        device = device)

for p in model.parameters():
    p.requires_grad = True

images = images.cpu().detach()
#images.shape

#visualize
images = images[0:16]

images = (images + 1) / 2

#create grid with size by nrow x ncol
img_grid = make_grid(tensor = images,
                     nrow = 8,
                     padding = 1,
                     pad_value = 1)

#conver to channel last
img_grid = img_grid.permute(1,2,0)

plt.figure(figsize=(16, 16))
plt.imshow(img_grid)
plt.axis('off')
plt.show()

```

```

[169]: generate_samples(model,
                      sampler,
                      iters = 50,
                      step_size = 10,
                      num_samples = 16,
                      image_size = (1, 28, 28),
                      device = device)

```



## 11 Complete training

```
[75]: def EBM_train(model,
    device,
    dataloader,
    optimizer,
    epochs,
    sampler,
    sampler_iters = 60,
    sampler_step_size = 10,
    alpha = 0.2, #regularization weight.
    scheduler = None,
    max_batches = None):

    #if model_path is not None and not os.path.exists(model_path):
    #    os.makedirs(model_path)

    loss_hist = []
    reg_hist = []
    cdiv_hist = []

    for epoch in range(1, epochs + 1):

        epoch_reg, epoch_cdiv, epoch_loss = EBM_train_one_epoch(
            model = model,
            device = device,
            dataloader = dataloader,
            optimizer = optimizer,
            epoch = epoch,
            sampler = sampler,
            sampler_iters = sampler_iters,
            sampler_step_size = sampler_step_size,
            alpha = alpha,
            max_batches = max_batches)

        if scheduler:
            #adjusting LR is necessary
            scheduler.step()

        loss_hist.append(epoch_loss)
        reg_hist.append(epoch_reg)
        cdiv_hist.append(epoch_cdiv)

    return reg_hist, cdiv_hist, loss_hist
```

```
[81]: reg_hist, cdiv_hist, loss_hist = EBM_train(model,
    device,
```

```

        dataloader,
        optimizer,
        epochs = 20,
        sampler = sampler,
        sampler_iters = 60,
        sampler_step_size = 10,
        alpha = 0.2,
        scheduler = scheduler,
        max_batches = None)

```

'Epoch 2 [127/469 (27%)]'

'\* regularization batch 0.000421 epoch 0.000440'

'\* contrastive divergence batch 0.020891 epoch 0.023432'

'\* loss batch 0.021313 epoch 0.023872'

KeyboardInterrupt

Traceback (most recent call last)

Cell In[81], line 1

```

----> 1 reg_hist, cdiv_hist, loss_hist = EBM_train(model,
        2         device,
        3         dataloader,
        4         optimizer,
        5         epochs = 20,
        6         sampler = sampler,
        7         sampler_iters = 60,
        8         sampler_step_size = 10,
        9         alpha = 0.2,
       10         scheduler = scheduler,
       11         max_batches = None)

```

Cell In[75], line 22, in EBM\_train(model, device, dataloader, optimizer, epochs, sampler, sampler\_iters, sampler\_step\_size, alpha, scheduler, max\_batches)

```

    18 cdiv_hist = []
    20 for epoch in range(1, epochs + 1):
----> 22     epoch_reg, epoch_cdiv, epoch_loss = EBM_train_one_epoch(
        23         model = model,
        24         device = device,
        25         dataloader = dataloader,
        26         optimizer = optimizer,
        27         epoch = epoch,
        28         sampler = sampler,
        29         sampler_iters = sampler_iters,
        30         sampler_step_size = sampler_step_size,
        31         alpha = alpha,
        32         max_batches = max_batches)
    34     if scheduler:

```



```

35         #adjusting LR is necessary
36         scheduler.step()

```

Cell In[56], line 50, in EBM\_train\_one\_epoch(model, device, dataloader, optimizer, epoch, sampler, sampler\_iters, sampler\_step\_size, alpha, max\_batches)

```

46 #noise = torch.randn_like(real_images) * math.sqrt(eta)
48 real_images.add_(noise).clamp_(min=-1.0, max=1.0)
---> 50 fake_images = sampler.train_sample(model = model,
49                                     iters = sampler_iters,
50                                     step_size = sampler_step_size,
51                                     device = device)
55 images = torch.cat([real_images, fake_images], dim = 0)
57 optimizer.zero_grad()

```

Cell In[54], line 45, in EBMSampler.train\_sample(self, model, num\_new, iters, step\_size, device)

```

38 # Perform MCMC sampling
39 images = self.langevine_dynamics_sample(model,
40                                     images,
41                                     iters = iters,
42                                     step_size = step_size,
43                                     device = device)
---> 45 self.buffer = list(images.to(torch.device("cpu")).chunk(self.
44                                     sample_size, dim=0)) + self.buffer
46 self.buffer = self.buffer[:self.buffer_len]
48 return images


```

KeyboardInterrupt:

```

[105]: generate_samples(model,
        sampler,
        iters = 80,
        step_size = 100,
        num_samples = 16,
        image_size = (1, 28, 28),
        device = device)

```



## 12 Image denoising with Langevine dynamics equation

```
[106]: def denoise_samples(model,
                        images,
                        sampler,
                        iters = 60,
                        step_size = 10,
                        num_samples = 16,
                        image_size = (1, 28, 28),
                        device = torch.device("cpu")):

    model.eval()
    model.to(device)

    for p in model.parameters():
        p.requires_grad = False

    images = sampler.langevine_dynamics_sample(
        model,
        images,
        iters = iters,
        step_size = step_size,
        return_images_per_step = False,
        device = device)

    for p in model.parameters():
        p.requires_grad = True

    images = images.cpu().detach()
    #images.shape

    #visualize
    images = images[0:16]

    images = (images + 1) / 2

    #create grid with size by nrow x ncol
    img_grid = make_grid(tensor = images,
                        nrow = 8,
                        padding = 1,
                        pad_value = 1)


    #conver to channel last
    img_grid = img_grid.permute(1,2,0)
```

```
plt.figure(figsize=(16, 16))
plt.imshow(img_grid)
plt.axis('off')
plt.show()
```

```
[131]: images, _ = next(iter(dataloader))
images = images[0:16]
images.shape
```

```
[131]: torch.Size([16, 1, 28, 28])
```


```
[133]: plot_images(images)
```




```
[140]: noise = torch.zeros_like(images)
noise.normal_(0, 1)
#noise = torch.rand(images.size())
noise = torch.randn(images.size())
images.data.add_(noise.data)
#noise_images = images + noise
noise_images = images.data.clamp_(min=-1.0, max=1.0)
noise_images.shape
```

```
[140]: torch.Size([16, 1, 28, 28])
```

```
[141]: plot_images(noise_images)
```



```
[142]: denoise_samples(model,  
                        noise_images,  
                        sampler,  
                        iters = 200,  
                        step_size = 10,  
                        num_samples = 16,  
                        image_size = (1, 28, 28),  
                        device = device)
```



```
[ ]:
```